

SVM

implement

- Fit a hyperplane that best classifies the data points while maximizing the margin, which is the distance between the nearest data point of each class and the hyperplane.
- Margin: distance between line and closest points
- Support vectors: points that determine the margin (i.e. lie on the margin, are closest together)
- note:
 - Points can be within the margin
 - Points can be on the wrong side of the hyperplane
 - Points on the right side of the margin have no bearing on the computation
 - Some points can be misclassified
 - Some points can be within the margin


Variations and Extensions

Kernel Functions: Transform data into higher dimensions to make it easier to separate with a hyperplane. Common kernels include linear, polynomial, and radial basis functions.

Parameters

C: “budget” for points on the wrong side of the margin

a larger C value will narrow the range of the minimum misclassification limit, while a smaller C value will expand the range, allowing more misclassified data.



γ Parameter: Applies in kernels like the Radial Basis Function (RBF).

Degree & coef0 : Applies in kernels like the polynomial.

Advantages and Disadvantages

Advantages:

Effective in high-dimensional spaces.

Still effective when the number of dimensions exceeds the number of samples.

Uses a subset of training points (support vectors), making it memory efficient.

Disadvantages:

Not suitable for large datasets as the training time with SVMs is high.

Less effective on noisier datasets with overlapping classes.

The choice of the kernel and regularization can have a large impact on performance.

Plots & Code:

The support vector classifier with two features can be visualized by plotting values of its *decision function*. We have included a function for this in the [ISLP](#) package (inspired by a similar example in the [sklearn](#) docs).

```
'''
```

```
from ISLP.svm import plot as plot_svm
fig, ax = subplots(figsize=(8,8))
plot_svm(X,
        y,
        svm_linear,
        ax=ax)
'''
```

```
'''
```

```
svm_linear = SVC(kernel='linear').fit(X, y)
svm_rbf = SVC(kernel='rbf').fit(X, y)
svm_poly = SVC(kernel='poly').fit(X, y)
'''
```

Cross validation

```
'''
```

```
kfold = skm.KFold(5,
                  random_state=0,
                  shuffle=True)
grid = skm.GridSearchCV(svm_rbf,
                        {'C':[0.1,1,10,100,1000],
                        'gamma':[0.5,1,2,3,4]},
                        refit=True,
                        cv=kfold,
                        scoring='accuracy');
grid.fit(X_train, y_train)
grid.best_params_
'''
```

ROC Curve and AUC Score: Useful for evaluating the classifier's ability to discriminate between classes.

```
'''
```

```
fig, ax = subplots(figsize=(8,8))
```

```
roc_curve(best_svm,
```

```
    X_train,
```

```
    y_train,
```

```
    name='Training',
```

```
    color='r',
```

```
    ax=ax);
```

```
'''
```